


TAKE-HOME WALKTHROUGH

PV FLEET EXPLORER

Triage 10,000 messy solar-site records —
see the fleet, trust the data, act on the gaps.

Debbie Chen · August 2026 · Next.js + React 19 · built with Claude Code

50 states



10,000 sites

Agenda

~15 minutes + questions — interrupt anytime.

I Understand

~3.5 min

- 01 Who this is for**
Maya, and the five questions she asks
- 02 The data - completeness**
10,000 records — missing a lot, rarely wrong
- 03 Field logic - validity**
Is the data reasonable? The data surfaced four kinds of suspicion

II The rules system

~4 min

- 04 Confidence levels**
before a rule gets power, it gets graded
- 05 Two systems - trust score & badges**
two places a rule can act - with very different blast radius
- 06 Rules mapping with two systems**
with different levels of power, where should each rule live?

III Show & decide

~8 min

- 07 Live demo**
system walkthrough
- 07A Map overview**
- 07B Site details**
- 07C Site compare**
- 07D Add site form**
- 08 Decisions**
five calls I would defend
- 09 Retrospective**
caught in my own review
- 10 Next steps**
from explorer to workbench
- 11 Closing**
what I'm really proposing is a framework.

I · UNDERSTAND


II · THE SYSTEM

III · LIVE DEMO & DECIDE

Understand

The user, the data — and the moment I caught my own map lying.

- Maya & her five questions
- The data: incomplete, not incorrect
- Field logic - validity



Maya reviews portfolios, not pixels.



Maya

Asset Data Analyst



Triage, not repair — find the 15% that will cause problems.



Knows solar cold — azimuth and tilt need no explaining.



Lives in Excel today — filters, sorting, yellow cells.



One afternoon per portfolio — not a week.

SHE ASKS FIVE QUESTIONS — THE APP ANSWERS IN ORDER

1

“What did we just get?”

the map

2

“Which ones are a problem?”

trust score & details

3

“Why is this one a problem?”

the breakdown

4

“Are these two the same?”

compare

5

“This one is missing.”

add site

HOW MIGHT WE

help a time-boxed reviewer explore 10,000 messy PV records, trust what she sees, spot what needs attention, and know what to do next?

10,000 records · 14 fields · missing a lot.

14 FIELDS → 5 GROUPS (the UI sections mirror these)

IDENTITY

system_ID

LOCATION

state

zip_code

latitude

longitude

INSTALL & OWNERSHIP

installation_date

third_party_owned

SCALE

system_size_DC

module_quantity_1

efficiency_1

PHYSICAL CONFIG

tracking

ground_mounted


azimuth_1

tilt_1

HOW “MISSING” ARRIVES

-1 · -1.0 · -9999 · “redacted” · blank — plus boolean-ish -1/0/1 and ZIPs stored as floats (“32605.0”). Normalized ONCE at the boundary; components only ever see number | null.

0 is a trap: azimuth 0 = north and tilt 0 = flat are VALID; size 0 is not.



7 fields

are required data for creating a site (ID, location, and system_size)

50% +

physical configuration data is missing

Is the data reasonable? The data surfaced four kinds of info/suspicion.



LOCATION

12.3% outside their state

ZIP, state, and coordinates cross-checked. The coords are privacy-fuzzed or mistakes?

- 180 < longitude < 180
- 90 < latitude < 90
- invalid zip
- invalid state
- 1,214 coords are out of state
- 140 unmapped kept
- 33 ZIP & state mismatch



INSTALLATION & OWNERSHIP

15 / 20 yr

Aging at 15, review at 20. Age is context —

- never lowers trust
- 653 within 15–19 yrs
- 56 exceed 20+ yrs
- 1710 third-party owned




SCALE

100–700 W / module [15-20%]

Size ÷ modules is the strongest cross-check. Efficiency is judged against its era.

- size > 0
- module (integer) > 0
- 659 implied W/module outside 100–700
- > 10,000 kW (extreme)
- 5 era cohorts (1998–2005 15.7% · 2006–2010 14.3% · 2011–2015 15.8% · 2016–2020 18.3% · 2021+ 20.6%)
- 0% < efficiency <= 100%
- 28 below 80% of cohorts mean's (low efficiency)
- <5% · >30% (extreme)



PHYSICAL CONFIG

5,432 unknown tracking

Missing angles only matter when tracking = fixed. Unknown is never assumed fixed.

- tracking → angles n/a ✓
- Azimuth 0-359
- Tilt 0-90
- 1,116 ground-mounted
- 481 in tracking system
- 200 tracking orientation reference only
- 2,156 fixed w/o angles
- 56 north-facing fixed

“ 24 candidate rules came out of these groups. Each one is a hypothesis; none has earned power yet.

I · UNDERSTAND

II · THE SYSTEM

III · LIVE DEMO & DECIDE

The *system*

I can only be sure of physics — so confidence decides how much power every rule gets.

- Confidence levels: physics · heuristic · open
- Two systems: score vs badges
- Rule placement: confidence × power



Before a rule gets power, it gets graded.

Physics

What I can verify alone — math and common sense.

-180 < longitude < 180

-90 < latitude < 90

invalid zip

invalid state

size > 0

module (integer) > 0

0% < efficiency <= 100%

Azimuth 0-359

Tilt 0-90

Heuristic

What data + AI experience suggest.
The numbers are mine — an expert should tune them.

653 within 15–19 yrs

56 exceed 20+ yrs

659 implied W/module outside 100–700

> 10,000 kW (extreme)

28 below 80% of cohorts mean's (low efficiency)

<5% · >30% (extreme)

Open questions

What only interviews can answer.
Suspicion without ground.

- Are the coordinates fuzzed or wrong?
- What information do reviewers care about?
- How to handle the physical rules for trust and info?

1,214 coords are out of state

140 unmapped kept

33 ZIP & state mismatch

1710 third-party owned

1,116 ground-mounted

481 in tracking system


200 tracking orientation reference only

2,156 fixed w/o angles


56 north-facing fixed

🔦 The data can tell you what is unusual. Only the field can tell you what is truly wrong. I built the framework so that conversation changes constants, not code.

Two places a rule can act — with very different blast radius.



Rule placement: confidence × power.



“ Under domain uncertainty, the system's job is to be useful without overclaiming — certainty scores, suspicion badges, hunches whisper.

BUILT TO BE CALIBRATED, NOT REBUILT

Thresholds, severities, and open questions are all first-pass judgments.
When domain experts weigh in, a confirmed rule unlocks the next cell and a debunked one retires — each move changes a named constant, never the logic.

I · UNDERSTAND

II · THE SYSTEM

III · LIVE DEMO & DECIDE


Live demo & decide

Five questions, five clicks — then the calls I made and what comes next.

- **Demo:** system walkthrough (map explore, site details, site compare, add site form)
- **Decisions & retrospective**
- **Next steps** · closing



Aggregate with states at the beginning.



One panel carries the whole argument.

SITE_08363

76

GOOD TRUST

Completeness 60 · Validity 100

Needs attention: 1

Aging system

LOCATION

StateColorado (CO)

ZIP80424

Coordinates38.8416, -104.9947

INSTALLATION & OWNERSHIP

Installed ⓘ2009-03-03

Age17 years

OwnershipUnknown

SCALE

System size2.6 kW DC

Modules —

Efficiency —

PHYSICAL CONFIGURATION

4

N

Tracking ⓘUnknown

MountUnknown

FacingSE · 148°

Tilt39°

NEEDS ATTENTION (1)

Tracking type is unknown — not assumed to be fixed, so orientation completeness cannot be fully assessed.

Add to comparison

Close

Why this score?

76

GOOD TRUST

completeness 60 × 0.6 + validity 100 × 0.4

Completeness 60

Validity 100

WHAT IS MISSING

60 / 100 applicable pts

✓ Coordinates20 / 20

✓ System size20 / 20

✓ Installation date10 / 10

✓ State + ZIP10 / 10

✗ Module quantity0 / 10

✗ Module efficiency0 / 10

✗ Tracking type0 / 10

✗ Mount type0 / 5

✗ Ownership0 / 5

Orientation (not applicable — tracking type unknown, fixed-tilt is not assumed)–

FILLING THESE WOULD RAISE THE SCORE

Module quantity+6 trust

Module efficiency+6 trust

Tracking type+6 trust

Mount type+3 trust

Ownership+3 trust

Tracking type is unknown — not assumed to be fixed, so orientation completeness cannot be fully assessed.
Flags never change the score.

Deterministic: the same inputs always produce the same score.

5

Efficiency

Module efficiency is missing.
Affects: Completeness –10 pts

Modules

Module count is missing — the implied module-power check cannot run.
Affects: Completeness –10 pts

- 1 Trust score — and its “why”**

The badge is a button. The popover shows the formula, the weighted checklist, and “+N pts if filled”. It renders checks[] and penalties[] verbatim — the UI cannot disagree with the math.
- 2 Three-severity badge system**


Red critical · amber needs-attention · blue info. Info never inflates the count — third-party ownership is a fact, not an alarm.
- 3 Four sections mirror the data**

LOCATION · INSTALLATION & OWNERSHIP · SCALE · PHYSICAL CONFIGURATION — the same five-group taxonomy from the CSV, so a reviewer always knows where a field lives. Missing renders as “—”, never blank.
- 4 2.5D physical configuration**

Tilt you can see, mount you can recognise, a yellow facing arrow — and a hard rule: tracking systems never get drawn with a fixed orientation. Exact degrees always sit beside it as text.
- 5 Flags where the field is**


■ sits next to the field it concerns. Click it: plain-language reason, the triggering number, the score impact, and the machine code that traces back to the spec.

When comparing, the sites should be highlighted.



Same site, three altitudes of comparison.

1 Small summary cards — the collapsed tray



2 Generated comparison summary


Site 1 is the largest selected system.

Only site 1 uses tracking.

Sites 3 & 4 are missing orientation data.

Max three pills, priority-ordered, computed only from derived fields — the tray reads the data so the reviewer starts with the differences that matter, not a wall of numbers.

3 Honest comparison viz side-by-side



A reviewer-friendly form, designed from messy data.

Add new PV site

Enter the known site details. Optional values can remain unknown.

1

LOCATION

State (required)

Texas (TX)

ZIP code (required)

78729

Look up ZIP

Austin, TX · centroid used

Latitude (required)

30.4521

Longitude (required)

-97.7688

Pick location on map

Coordinates from ZIP centroid — approximate to the ZIP area, not the exact site. Pick on the map or enter exact values to refine.

Installation date (required)

2026/08/20

▲ Installation date cannot be in the future.

SYSTEM

System size (kW DC) (required)

1000000

Enter the DC system capacity in kilowatts.

Module quantity

20

Module efficiency (%)

20.5

4

PHYSICAL CONFIGURATION

Tracking type

Unknown

Fixed

Tracking

Orientation can be entered, but completeness cannot be fully assessed until tracking type is known.

Mounting type

Unknown

Roof

Ground

Azimuth (0–359°)

180

Tilt (0–90°)

25

0° N · 90° E · 180° S · 270° W

OWNERSHIP

Ownership

Unknown

Owner-held

Third-party

▲ Worth confirming — these do not block submission: This is unusually large for the sample fleet. Please confirm.

2

SITE_10001 will be generated

Cancel

Add site

1 Required fields

Without these it cannot be a site. (state, zip_code, latitude, longitude, installation date, system size)

2 Generate ID by system

To ensure uniqueness, the system generates the ID automatically.

3 Look up ZIP / Pick location on map for user-friendly choices

For UX and correctness, we provide alternative way for reviewers to easily add site location info.

4 Optional and unknown might be the real-world answer

I hesitate how hard we should control the form, but the messy data told us this might be the answer.

5 Hard gate vs soft warning

Blocking a plausible-but-odd value would force users to lie to the form.

Errors are for the impossible; everything merely suspicious becomes a flag on the record.

BLOCKS ▲


- missing required field
- azimuth 360+ · tilt 91+
- efficiency > 100% · size ≤ 0
- module count 1.5 (not whole)
- future installation date
- ZIP not in the US database

WARNS — still submits

- fixed system, no angles → FIXED_ORIENTATION_MISSING
- efficiency < 5% or > 30%
- implied W/module $\notin$ 100–700
- 20+ year installation
- ZIP belongs to another state
- coords far from the ZIP


✓ Submit runs the same pipeline as the CSV import — validate → generate ID → normalize → score → flag. The site you typed lands on the map as a first-class record: scored, flagged, comparable, indistinguishable from an imported one.

Five calls I would defend.



Nulls decided once

Sentinels die at one boundary — below it, everything is number | null.




The map IS the basic filter

Geography is the reviewer's natural first cut — so state aggregation ships as the working filter: click a state, land in its sites.



Make every detail explain itself

The record level got the depth: a score breakdown that renders the math verbatim, the 2.5D physical configuration you recognize instead of decode, and every badge carrying its reason, its trigger, and its score impact.




Ship rules with confidence

The score judges the record. Badges describe the site. Both can be true at once.



Weights were measured

0.6 / 0.4 chosen against the real fleet — validity averages 99, so completeness must lead.



Strict door, gentle door

The form blocks a human typing nonsense. Imports get flags, not 1,214 alarms.

Caught in my own review.

	PROCESS	Figma changed, code didn't	fix one spec source + visual check per phase
	EVIDENCE	Badge design is a guess	fix A/B it — measure time-to-find, not taste
	CRAFT	Raw decimals leak through	fix format once, at the display boundary
	CRAFT	Northeast bubbles collide	fix ~30 lines of pairwise repulsion
	CRAFT	Markers vs loud basemap	fix the desaturate filter never applied — one CSS rule

→ None of these blocked the reviewer — which is exactly why they were easy to miss and worth writing down.

A better way to handle location mismatches

10,000 sites

8,646 mapped

1,354 need location review

→

persistent banner

· queue is one click away

Location review

1,354 sites

Conflicting sources

1,214

No coordinates

140

🚩 SITE_00002

3.84 kW · installed 2013

DECLARED · STATE + ZIP

Texas · ZIP 79830

Alpine, TX

centroid 30.263, -103.654

RECORDED · COORDINATES

29.105, -101.416

falls outside Texas

(point-in-state check)

sources ≈ 250 km apart



Keep coordinates

trust the recorded point

Use ZIP centroid

approximate · marked on map

Send back

ask the source

Every choice is recorded: resolved by · source used · timestamp — the record keeps its history.

🚩 SITE_00049

VT · 72433

78 km outside VT

▼

🚩 SITE_00100

AR · 71653

across the MS border

▼


🚩 SITE_00170

VA · 24970

122 km from ZIP

▼

From explorer to workbench.



 All are already designed — with real counts. 1-3 for next steps. 4 is the horizon.

What I'm really proposing is a framework.

Trust scores, flags, and review queues — with every threshold as a named, tunable dial. The structure is mine; the calibration belongs to the field.

Measure before deciding


Explainable by construction

Conversation changes constants, not code

A lot of what you saw are first-pass judgments — the thresholds, the severities, the open questions. They're built to be corrected by people who know the field better than I do.


**So — I'm very open to your feedback, and genuinely: this was a joy to think through.
I learned a lot building it, and I'd love to hear where you'd push back.**

Thank you.



APPENDIX

APPENDIX



Trust score details

COMPLETENESS · WEIGHTS

Coordinates (pair · half credit)	20
System size	20
Installation date	10
State + ZIP	10
Module quantity	10
Module efficiency	10
Tracking type	10
Mount type	5
Ownership	5
Orientation (pair) — fixed only	15

denominator = 115 fixed · 100 tracking/unknown
Tracking & unknown: orientation leaves the denominator — never counted as missing.

VALIDITY · PENALTIES (supplied values only)

OUT OF RANGE	
lat/lng outside the globe	-30
system size ≤ 0	-30
azimuth ∉ 0–359 · tilt ∉ 0–90	-20
efficiency ∉ 0–100%	-20
module qty not whole ≥ 1	-15
unreadable / bad ZIP format	-10
IMPLAUSIBLE	
implied W/module ∉ 100–700	-15
efficiency <5% or >30%	-15
size > 10,000 kW	-10

Fleet reality: 659 / 10,000 lose any points — all to implied module power.

WORKED EXAMPLE · SITE_00002 (REAL)

8 of 10 applicable checks present	C = 80
missing: tracking · modules	-20 pts
orientation	excluded
nothing supplied is invalid	V = 100

0.6 · 80 + 0.4 · 100 = 48 + 40 = 88 · HIGH


GUARANTEES

- checks[] and penalties[] reconcile with the score — tested over all 10,000 records
- deterministic: same inputs, same score
- age, ownership, mounting, tracking never lower it
- one constant to retune the 0.6/0.4 weighting

Badges details



AI wrote the code. The judgment stayed manual.



- Specs went in as commit #1
- One phase per prompt, /clear between
- Every diff read before commit
- Anything I can't explain doesn't ship

The folder structure is the argument.

Judgment lives in /utils as pure functions; /components only render what was already decided.

src/utils — PURE LOGIC

normalize.ts

validate.ts

metrics.ts

fieldStatus.ts

geo.ts

aggregate.ts

compare.ts

ids.ts

states.ts

__tests__/

sentinels die here → number | null

form + import rules · errors vs warnings

trust = 0.6·C + 0.4·V, with checks[]

the badge engine · critical / review / info

ZIP · state · coords cross-checks

state bubbles · counts & capacity

tray slots · stable numbering · insights

SITE_XXXXX generation — never typed

state names, centroids, lookups

140 tests — every rule pinned

Zero React imports in this folder. Every rule runs in a test, not a hook.

src/components — THIN VIEWS

FleetExplorer.tsx

FleetMap.tsx + StateBubbleLayer

SummaryBar.tsx

SiteDetailPanel.tsx

TrustRing.tsx + ScoreBreakdownPopover

FieldStatusButton.tsx

OrientationDiagram.tsx

CompareTray.tsx

AddSiteDrawer.tsx

reducedMotion.ts

the shell — owns selection state

aggregation → clusters → pins

honest fleet counts

four sections · mirrors the CSV taxonomy

renders checks[] verbatim

the 🚩 next to the field it concerns

the 2.5D config — recognition, not decoding

2–5 sites · numbers never shift

same pipeline as the CSV import

a11y — motion respects the OS

Components ask utils; they never decide. If a number is wrong, the fix is in one folder.

APPENDIX 05 MAP FILTER — BASIC / ADVANCED

Overlay, top-left of the map. Every option shows its real count from pv_sites_sample.csv so the reviewer knows what a filter costs before applying it.

Filter10,000 shown

State, ZIP, or site ID

Enter a ZIP or state to zoom there

TRUST

All10,000

High3,723

Good6,047

Limited230

NEEDS ATTENTION

Any review flag6,411

Missing coordinates140

Advanced filters+7

Filter847 of 10,000

CAFixed-tiltSouth-facing10-50 kW

Clear all

HARDWARE · MOUNTING

Any10,000

Roof2,662

Ground1,116

Unknown6,222

HARDWARE · SYSTEM TYPE

Any10,000

Fixed4,087

Tracking481

Unknown5,432

ORIENTATION · FACING

South1,987

East355

West428

North460

Unknown6,770

Only 3,230 sites record an azimuth. Filtering by facing hides the rest.

ORIENTATION · TILT

0°

15° - 40°

90°

SCALE · SYSTEM SIZE

< 5 kW2,103

5 – 10 kW4,749

10 – 50 kW2,620

50 – 500 kW250

500 kW +104

INSTALLATION AGE

Any10,000

2021+4,108

2016–203,825

≤20152,067

FLAGS

Tracking type unknown5,432

Fixed missing orientation2,156

Unusual implied module power659

Older installation (20+ yr)56

North-oriented fixed array56

Third-party owned1,710

Ground-mounted1,116

B – Advanced — same panel, expanded. Scrols in context, max-height 70vh.

SEARCH → ZOOM BEHAVIOUR

787

ZIP 78701Austin, TX1 site

ZIP 787...12 matching ZIP prefixes–

Texasstate783 sites

SITE_00787TX · 6.2 kWjump

· ZIP (5 digits) → fly to that ZIP centroid at zoom 11

· ZIP prefix (3–4 digits) → fit bounds of all matching sites

· State code or name → fit that state, land at ≥ zoom 6

· SITE_XXXXX → fly to it and open its detail panel

· No match → inline "No sites match 999xx", map unchanged

Search moves the camera; it does not filter. Filtering hides records — a reviewer searching for one ZIP still wants the neighbours visible for context.

BEHAVIOUR RULES

COUNTS ARE LIVE

Every option shows how many sites it would leave. Counts recompute against the OTHER active filters, so a zero option is visibly a dead end before the reviewer clicks it.

THE MAP OBEYS THE FILTER

State bubbles and clusters both recount. The summary bar shows "847 of 10,000 shown". Filtering never changes the trust score of anything — it only changes what is on screen.

FILTER ≠ SEARCH

Filter narrows the working set. Search moves the camera. Keeping them separate avoids the classic bug where typing a ZIP silently deletes the rest of the fleet.

UNKNOWN IS SELECTABLE

"Unknown" is a first-class option on every tri-state facet — 6,222 sites have no mounting type, and finding them is a real review task, not an edge case.

ACTIVE FILTERS STAY VISIBLE

Chips at the top of the panel, each removable, plus one "Clear all". When the panel is collapsed the chips remain, so a filter can never be silently on.

EMPTY RESULT IS A STATE

"0 of 10,000 match — clear filters" with a working action, not a blank map.

IF TIME IS SHORT — BUILD IN THIS ORDER

1Trust segmented control + live count in the summary bar8m

2Search box → state / ZIP / site ID camera moves10m

3Needs-attention checkboxes (reuses the flag data)6m

4Advanced disclosure: mounting, system type, size buckets12m

5Facing chips + tilt range slider10m

6Per-option live counts recomputed against other filters10m

Steps 1–3 are the whole value: they turn 10,000 records into a working set. Everything below is polish — stop wherever the clock says stop and note the rest in DECISIONS.md.

APPENDIX 06 LOCATION REVIEW — QUARANTINE & RESOLVE

Conflicted locations leave the map and enter a queue. The reviewer — not the renderer — decides which source to trust.

10,000 sites8,646 mapped1,354 need location review →

persistent banner · queue is one click away

Location review1,354 sites

Conflicting sources1,214No coordinates140

SITE_000023.84 kW · installed 2013


DECLARED · STATE + ZIP

Texas · ZIP 79830
Alpine, TX
centroid 30.263, -103.654

RECORDED · COORDINATES

29.105, -101.416
falls outside Texas
(point-in-state check)

sources ≈ 250 km apart



Keep coordinates
trust the recorded point

Use ZIP centroid
approximate · marked on map

Send back
ask the source

Every choice is recorded: resolved by · source used · timestamp — the record keeps its history.

SITE_00049VT · 7243378 km outside VT

SITE_00100AR · 71653across the MS border

SITE_00170VA · 24970122 km from ZIP

MARKER VOCABULARY

Verified coordinates

reviewer kept the recorded point

ZIP centroid · approximate

dashed ring = position is the ZIP area, not the site

Unresolved — not on the map

in the review queue; never silently placed

THE FLOW

1Import detects a conflict → site is quarantined, not plotted

2Banner counts it: “1,354 need location review”

3Reviewer opens the queue, sees both claims + the distance

4Picks a source (or sends back) — choice is recorded

5Site joins the map in its earned marker state

Nothing on the map is unverified. Nothing unverified is invisible.

Score impact: unresolved location keeps the completeness credit for the fields that exist — quarantine is a workflow state, not a data penalty. Resolving to centroid keeps an “approximate” provenance flag.

APPENDIX 07 EDIT & DELETE — WITH PROVENANCE

The audit trail IS the feature. A record can change and stay auditable — because every change carries who, what, when, and why.

SITE_00121

EDITING

Esc cancels

SCALE

System size4.70 kW DC

Modules — →14

Source *

Installer email 08/04 ▾

Note

confirmed by installer

Edits run the same validation as the add-site form — an azimuth of 400 is just as impossible here.

🔗

Score preview: trust 87 → 92 · completeness +10 pts · UNUSUAL_IMPLIED_MODULE_POWER will re-run (4,700 W ÷ 14 = 336 W ✓)

Save change

Cancel

HISTORY

3 events

maya@

modules — → 14 · source: installer email · trust 87 → 92

2026-08-06 14:12

maya@

location resolved → ZIP centroid (approximate) · source: review queue

2026-08-06 11:03

import

created from pv_sites_sample.csv row 121

2026-08-05 20:31

Nothing is overwritten — an edit appends. The current value is the top of a stack you can always read down.

DELETE IS RETIRE, NOT ERASE

Retire SITE_00121?

The site leaves the map and every aggregate, but the record and its history remain readable. The ID is never reused.

Reason *

Duplicate of SITE_00119 ▾

Retire site

Cancel

Afterwards:

SITE_00121

RETIRED

searchable · history intact · restorable

FOUR RULES

1

An edit is an event, not an overwrite — source and author are required fields, like any other value.

2

Edits revalidate and rescore live: the reviewer sees the trust impact before saving.

3

Delete retires. Aggregates and the map forget the site; the record never forgets itself.

4

Manually-set values win over re-imports — and carry a “human-verified” provenance mark the score can trust.

APPENDIX 08 REVIEWER WORKLIST — TRIAGE BECOMES THROUGHPUT

Every review flag becomes a to-do with progress state. The map is one view of the workflow — no longer the workflow itself.

This week's review

214 of 1,522 done

North-oriented fixed arrays

NORTH_ORIENTED_FIXED_ARRAY

12 of 56 checked

SITE_04217

TX · 210 kW

az 352° · dev 172°

Open

Skip

SITE_00871

MN · 44 kW

az 8° · dev 172°

Open

Skip

SITE_02203

VT · 12 kW

az 341° · confirmed 2h ago

azimuth confirmed

Sorted by system size — biggest impact first

next unchecked

Location review

COORDS_CONFLICT + MISSING

146 of 1354 checked

Unusual implied module power

UNUSUAL_IMPLIED_MODULE_POWER

56 of 659 checked

Low efficiency for its era

LOW_EFFICIENCY_FOR_ERA

0 of 28 checked

THE MAP BECOMES ONE VIEW

Showing: this worklist (44)

Opening a worklist filters the map to its sites; everything else dims. Same fleet, new lens — “where are my 44 remaining north-facing arrays?” is one click.

WHY THIS IS THE HORIZON

1

Checking an item is itself an audited event — the worklist writes the same history as an edit.

2

Progress persists per flag type: “12 of 56” survives the weekend. Triage stops restarting from zero.

3

Done means decided, not deleted — confirming “yes, it really faces north” clears the to-do and keeps the flag as context.

4

Every count is the flag engine's output — the worklist adds no new judgment, only sequencing.